

Scale out,
a mówili, że się nie da

Arkadiusz Borek
LJUG

JDD 2016



@arekborek



arek.borek@gmail.com



github.com/arekborek



arekborek.blogspot.com/



Microservices?

You probably heard a lot already!

AGENDA

So many services

Deploy


Manage

Discovery

How?


Old way: shared machines

- No isolation
- No namespaces
- Common libs
- High coupled apps and OS




Old way: shared app servers

- No isolation
- No namespaces
- Common libs




Old way: many virtual machines

- Some isolation
- Expensive and inefficient
- Still highly coupled to the guest OS
- Hard to manage



New way: containers



Containers: Process isolation

- CPU (PIDs)
- Memory
- Network interfaces
- Filesystem

Containers: Security

- Chroots determines what parts of the filesystem a user can see
- Control - CPU, Memory, IO via cgroups
- Limits what a user can do: mount, kill, chown
- No access to host !!!

Containers: Other features

- New approach to packaging
- Lightweight
- Runs on any machine
Physical, virtual, cloud
- Write once – run everywhere!
Development, Test, Staging, Production



docker

Let's Containerize

Containerize Option #1

Docker file

Containerize Option #2
Maven plug-in made by spotify

Containerize Option #3
Docker hub / GitHub

Example of docker

Docker file

/hellonode

- Dockerfile
- server.js

server.js

```
var http = require('http');  
var handleRequest = function(request, response) {  
  console.log('Received request for URL: ' + request.url);  
  response.writeHead(200);  
  response.end('Hello World (' + process.env.HOSTNAME + ':V1) !');  
};  
var www = http.createServer(handleRequest);  
www.listen(8080);
```


Dockerfile

```
FROM node:4.4  
EXPOSE 8080  
COPY server.js .  
CMD node server.js
```

Build Docker Container

```
docker build \  
-t gcr.io/$PROJECT_ID/hello-jdd:v1 \  
.
```

Run Docker Container

```
docker run \  
-d \  
-p 8080:8080 \  
--name hello-jdd \  
gcr.io/$PROJECT_ID/hello-jdd:v1
```


Test

```
for ((;;)) do  
  curl http://localhost:8080;  
  printf '\n';  
  sleep 1;  
done
```

What next ??

- What if there are more containers?
- Scheduling – where should my containers run ?
- Lifecycle and health: keep my containers running
- Discovery: where are my container now ?
- Monitoring: What's happening with my container ?
- Scaling: Making job bigger or smaller

How to handle it ?

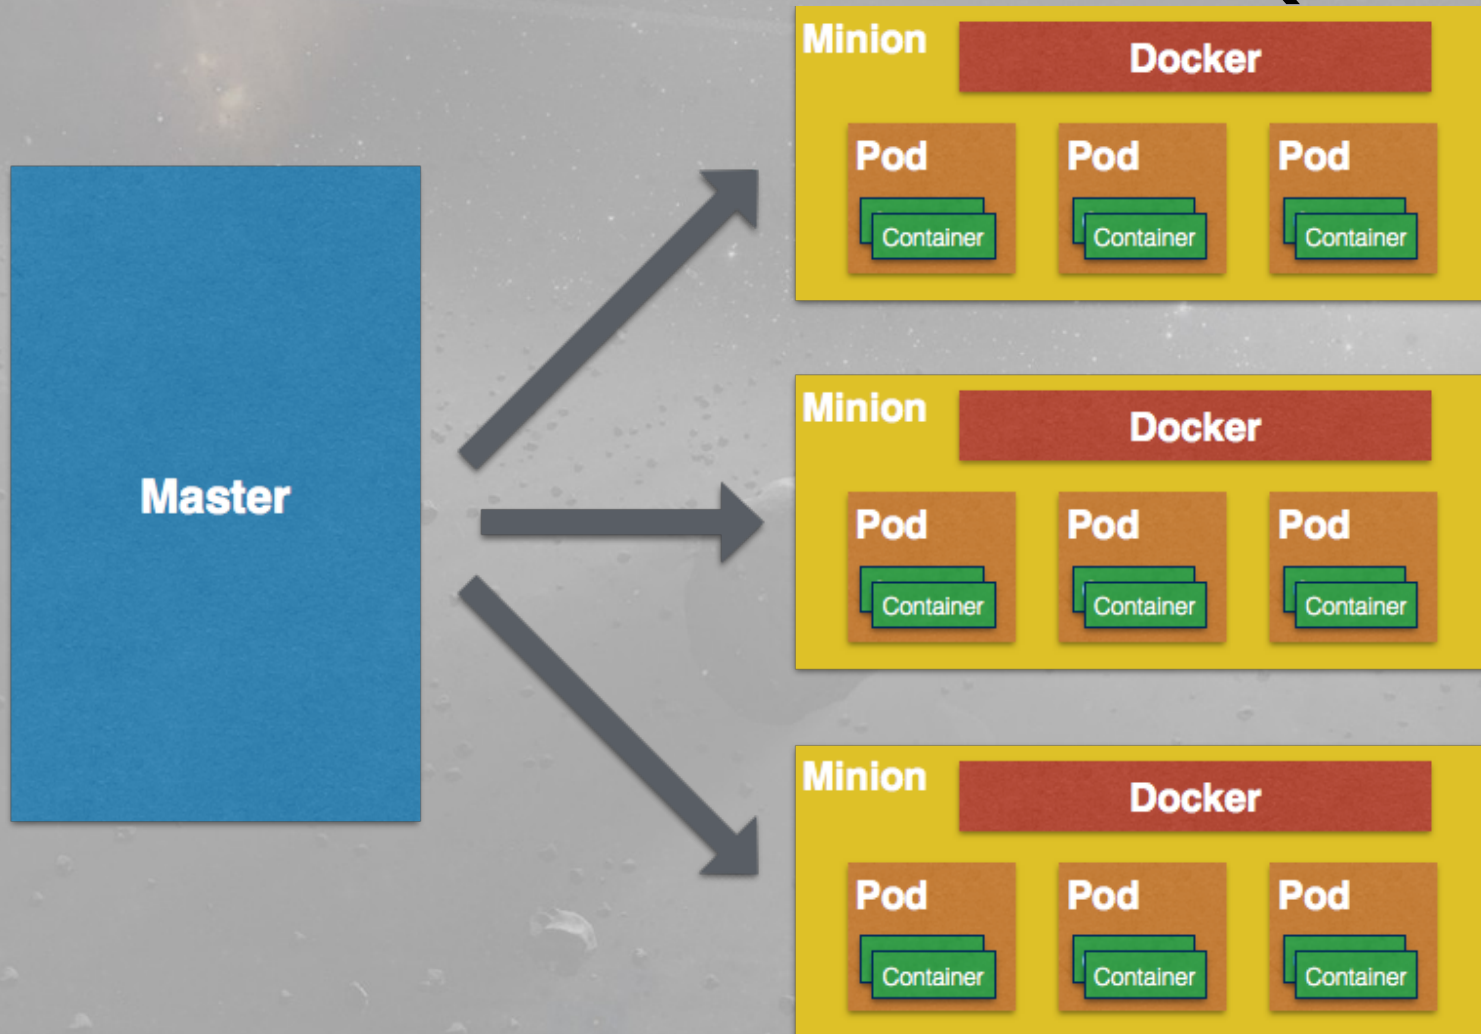


Let's kubectl-ing world!

Enter Kubernetes (k8s)

- Greek for “Helmsman”, also the root of the word “Governor”
- Container orchestrator, run containers
- Support multiple cloud and bare-metal environments
- Inspired and informed by Google’s experiences and internal systems
- Open source, written in Go

Enter Kubernetes (k8s)



Enter Kubernetes (k8s)

- Kubernetes master – order others
- Kubelet – ask how many resources you have if enough run machine
- Distributed process scheduler
- Cluster of machines as one
- All your machines via kubernetes can be treated as one single machine

k8s basis concepts

- Pod

Pod

- Group of containers
- One atomic resource that can be managed by Kubernetes
- Live and die together
- Shared network interface

k8s basis concepts

- Pod
- **Pod networking**


Pod networking

- Pod IPs routable, docker default is private IP
- Pods can reach each other without NAT, even across nodes
- No brokering of port numbers
- This is fundamental requirement, several SDN solutions

k8s basis concepts

- Pod
- Pod networking
- **Replication Controllers**

Replication Controller



- Allow start pods multiple times
- Check state if one of you pod died, RC run next instance of container for you

k8s basis concepts

- Pod
- Pod networking
- Replication Controllers
- **Service**

Service

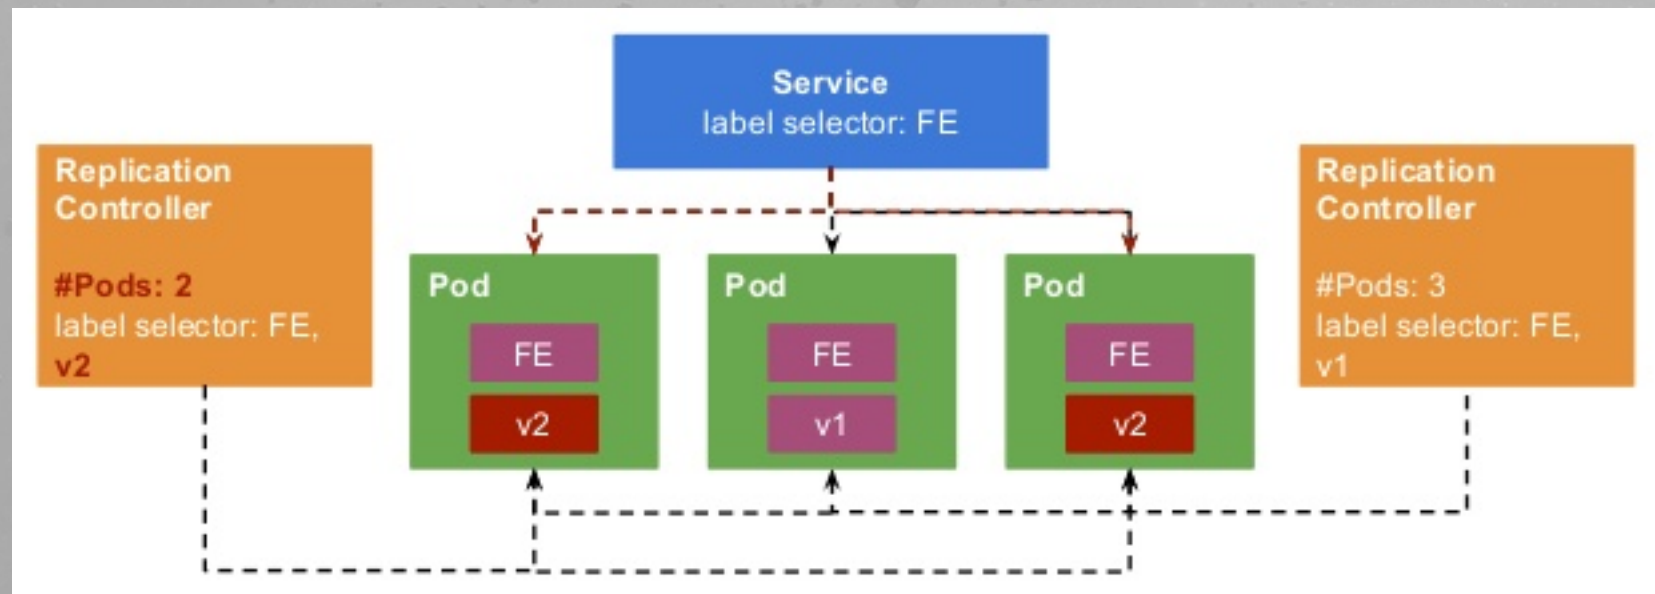
- Provide access to your N pods
- LB – front of pods
- Only route traffic to pods
- Selector-label concept

k8s basis concepts

- Pod
- Pod networking
- Replication Controllers
- Service
- **Labels & Selectors**

Labels & Selectors

- Label anything
- Name-value pair
- Make your own



Example of kube

Push image

```
gcloud docker push \  
gcr.io/$PROJECT_ID/hello-jdd:v1
```


Replication Controller

```
kubectl create -f rc-v1.yaml
```

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: hello-jdd-v1
spec:
  replicas: 1
  template:
    metadata:
      name: hello-jdd-v1
      labels:
        app: hello-jdd-v1
        env: jdd
        tier: backend
        version: v1
        visualize: "true"
    spec:
      containers:
        - name: hello-jdd
          image: "gcr.io/strong-art-145220/hello-jdd:v1"
          imagePullPolicy: Always
          ports:
            - containerPort: 8080
```

Service

Allow external traffic

```
kubectl create -f service.yaml
```



```
apiVersion: v1
kind: Service
metadata:
  labels:
    env: jdd
    visualize: "true"
  name: hello-jdd
spec:
  ports:
    - port: 8080 #The port that this service should serve on.
  selector: #Label keys and values connect RC and Service
    env: jdd
    tier: backend
  type: LoadBalancer
```


Test :)

```
for ((;;)) do \  
curl http://localhost:8080; \  
printf '\n'; \  
sleep 1; \  
done
```

Scale out

```
kubectl scale rc hello-jdd-v1 --replicas=2
```

Roll out an upgrade

```
gcloud docker push \  
gcr.io/$PROJECT_ID/hello-jdd:v2
```

```
kubectl create -f rc-v2.yaml  
kubectl scale rc hello-jdd-v1 --replicas=1  
kubectl scale rc hello-jdd-v2 --replicas=2  
kubectl scale rc hello-jdd-v1 --replicas=0
```


--rollback

```
kubectl scale rc hello-jdd-v1 --replicas=1  
kubectl scale rc hello-jdd-v2 --replicas=1  
kubectl scale rc hello-jdd-v1 --replicas=2  
kubectl scale rc hello-jdd-v2 --replicas=0
```


Some useful cmds

kubectl get rc

kubectl get pods

kubectl logs <POD_NAME>

kubectl cluster-info

kubectl config view

kubectl describe rc <RC_NAME>

kubectl get -o json pod

Some useful ...

Staging vs. production

Use label – deploy in the same infrastructure

Service discovery

Kubernetes API Or... DNS Lookups!

Use Docker Machine

In the cloud – faster network to download images

Some useful ...

Don't Log to Container Filesystem!

Log to a volume... `docker -v /tmp/log:/log` Or
Send it elsewhere!

Clean up disk spaces

Every image, layer and even container litters

Use Spring Profile

One container – run on Docker Machine,
Kubernetes, and App Engine!

The background of the slide is a grayscale image of a dense field of asteroids or debris in space. A large, irregularly shaped asteroid is prominent in the lower right foreground, showing a cratered surface. Numerous smaller, spherical and irregular objects are scattered throughout the mid-ground and background, creating a sense of depth and vastness. The overall tone is somber and dramatic, fitting the 'THE END !!' text.

THE END !!